

In [1]:

```
!pip install dmba
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)

Requirement already satisfied: dmba in /usr/local/lib/python3.8/dist-packages (0.1.0)

In [1]:

```
%matplotlib inline
import math
from pathlib import Path
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import ElasticNet
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import ExtraTreesRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.metrics import accuracy_score, roc_curve, auc
from sklearn.metrics import confusion_matrix
from dmba import regressionSummary, classificationSummary

from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import BayesianRidge
import matplotlib.pyplot as plt

custom_params = {"axes.spines.right": False, "axes.spines.top": False}
sns.set_theme(style="white", palette="Set2", rc=custom_params)
```

In [2]:

```
area_df=pd.read_csv("fulfilment_center_info.csv")
area_df.head()
```

Out[2]:


|   | center_id | city_code | region_code | center_type | op_area |
|---|-----------|-----------|-------------|-------------|---------|
| 0 | 11        | 679       | 56          | TYPE_A      | 3.7     |
| 1 | 13        | 590       | 56          | TYPE_B      | 6.7     |
| 2 | 124       | 590       | 56          | TYPE_C      | 4.0     |
| 3 | 66        | 648       | 34          | TYPE_A      | 4.1     |
| 4 | 94        | 632       | 34          | TYPE_C      | 3.6     |

In [3]:

```
area_df['center_type'].value_counts().plot(kind='barh',color="red")
```

Out[3]:

<AxesSubplot:>



In [4]:

```
print("Number of different centers from where the order dispatched")  
len(area_df['center_id'].unique())
```

Number of different centers from where the order dispatched

Out[4]:

77

In [5]:

```
meal_df=pd.read_csv("meal_info.csv")  
meal_df.head(7)
```

Out[5]:

|   | meal_id | category  | cuisine |
|---|---------|-----------|---------|
| 0 | 1885    | Beverages | Thai    |
| 1 | 1993    | Beverages | Thai    |
| 2 | 2539    | Beverages | Thai    |
| 3 | 1248    | Beverages | Indian  |
| 4 | 2631    | Beverages | Indian  |
| 5 | 1311    | Extras    | Thai    |
| 6 | 1062    | Beverages | Italian |


In [6]:

```
meal_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 51 entries, 0 to 50  
Data columns (total 3 columns):  
 #   Column      Non-Null Count  Dtype    
---  -  
 0   meal_id     51 non-null     int64    
 1   category    51 non-null     object    
 2   cuisine     51 non-null     object    
dtypes: int64(1), object(2)  
memory usage: 1.3+ KB
```

In [7]:

```
meal_df['cuisine'].value_counts().plot(kind='pie', autopct='%0.0f%%', startangle=90, wedgepro  
my_circle=plt.Circle( (0,0), 0.70, color='white')  
p=plt.gcf()  
p.gca().add_artist(my_circle)  
plt.show()
```




In [8]:

```
meal_df['category'].value_counts().plot(kind='bar',color="blue")
```

Out[8]:

<AxesSubplot:>



In [9]:

```
df=pd.read_csv("Food demand.csv")
df
```

Out[9]:

|      | id      | week | center_id | meal_id | checkout_price | base_price | emailer_for_promotion | h   |
|------|---------|------|-----------|---------|----------------|------------|-----------------------|-----|
| 0    | 1000000 | 3    | 157       | 2760    | 233.83         | 231.83     |                       | 0   |
| 1    | 1000001 | 100  | 104       | 2956    | 486.03         | 583.03     |                       | 0   |
| 2    | 1000002 | 143  | 75        | 1971    | 328.86         | 327.86     |                       | 0   |
| 3    | 1000003 | 41   | 24        | 2539    | 145.53         | 145.53     |                       | 0   |
| 4    | 1000004 | 45   | 83        | 2539    | 95.06          | 120.34     |                       | 0   |
| ...  | ...     | ...  | ...       | ...     | ...            | ...        |                       | ... |
| 1994 | 1002177 | 89   | 72        | 1311    | 130.04         | 177.51     |                       | 0   |
| 1995 | 1002178 | 24   | 50        | 2444    | 604.31         | 606.31     |                       | 0   |
| 1996 | 1002179 | 43   | 88        | 1971    | 291.06         | 291.06     |                       | 0   |
| 1997 | 1002180 | 107  | 58        | 1543    | 473.39         | 473.39     |                       | 0   |
| 1998 | 1002181 | 105  | 177       | 2322    | 284.27         | 284.27     |                       | 0   |

1999 rows × 9 columns

In [10]:

```
df.describe()
```

Out[10]:

|       | id           | week        | center_id   | meal_id     | checkout_price | base_price  | emailer_for_promotion |
|-------|--------------|-------------|-------------|-------------|----------------|-------------|-----------------------|
| count | 1.999000e+03 | 1999.000000 | 1999.000000 | 1999.000000 | 1999.000000    | 1999.000000 |                       |
| mean  | 1.001093e+06 | 75.393197   | 81.649825   | 2010.123562 | 327.302596     | 347.972866  |                       |
| std   | 6.323493e+02 | 41.743802   | 46.139173   | 554.686525  | 150.906902     | 158.625091  |                       |
| min   | 1.000000e+06 | 1.000000    | 10.000000   | 1062.000000 | 65.020000      | 93.120000   |                       |
| 25%   | 1.000546e+06 | 40.000000   | 43.000000   | 1543.000000 | 222.645000     | 242.530000  |                       |
| 50%   | 1.001094e+06 | 78.000000   | 76.000000   | 1971.000000 | 292.030000     | 309.430000  |                       |
| 75%   | 1.001638e+06 | 111.500000  | 110.000000  | 2539.000000 | 435.530000     | 447.230000  |                       |
| max   | 1.002181e+06 | 145.000000  | 186.000000  | 2956.000000 | 767.330000     | 767.330000  |                       |

In [11]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1999 entries, 0 to 1998
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    1999 non-null   int64
 1   week                  1999 non-null   int64
 2   center_id             1999 non-null   int64
 3   meal_id               1999 non-null   int64
 4   checkout_price        1999 non-null   float64
 5   base_price            1999 non-null   float64
 6   emailer_for_promotion 1999 non-null   int64
 7   homepage_featured     1999 non-null   int64
 8   num_orders            1999 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 140.7 KB
```

In [12]:

df.isnull().sum()

Out[12]:

```
id                0
week              0
center_id         0
meal_id           0
checkout_price    0
base_price        0
emailer_for_promotion 0
homepage_featured 0
num_orders        0
dtype: int64
```

In [13]:

df=df.drop(columns=['emailer\_for\_promotion','homepage\_featured'])

In [14]:

df.head()

Out[14]:

|   | id      | week | center_id | meal_id | checkout_price | base_price | num_orders |
|---|---------|------|-----------|---------|----------------|------------|------------|
| 0 | 1000000 | 3    | 157       | 2760    | 233.83         | 231.83     | 149        |
| 1 | 1000001 | 100  | 104       | 2956    | 486.03         | 583.03     | 161        |
| 2 | 1000002 | 143  | 75        | 1971    | 328.86         | 327.86     | 149        |
| 3 | 1000003 | 41   | 24        | 2539    | 145.53         | 145.53     | 540        |
| 4 | 1000004 | 45   | 83        | 2539    | 95.06          | 120.34     | 271        |

In [15]:

```
df=pd.merge(df,area_df,on='center_id')
```

In [16]:

```
df.head()
```

Out[16]:

|   | id      | week | center_id | meal_id | checkout_price | base_price | num_orders | city_code | regi |
|---|---------|------|-----------|---------|----------------|------------|------------|-----------|------|
| 0 | 1000000 | 3    | 157       | 2760    | 233.83         | 231.83     | 149        | 609       |      |
| 1 | 1000251 | 126  | 157       | 2306    | 338.53         | 340.53     | 15         | 609       |      |
| 2 | 1000336 | 16   | 157       | 2492    | 445.23         | 447.23     | 55         | 609       |      |
| 3 | 1000406 | 103  | 157       | 1109    | 192.09         | 339.50     | 68         | 609       |      |
| 4 | 1000520 | 128  | 157       | 1230    | 363.78         | 363.78     | 190        | 609       |      |

In [17]:

```
df=pd.merge(df,meal_df,on='meal_id')
```

In [18]:

```
df.head()
```

Out[18]:

|   | id      | week | center_id | meal_id | checkout_price | base_price | num_orders | city_code | regi |
|---|---------|------|-----------|---------|----------------|------------|------------|-----------|------|
| 0 | 1000000 | 3    | 157       | 2760    | 233.83         | 231.83     | 149        | 609       |      |
| 1 | 1001086 | 144  | 157       | 2760    | 184.36         | 261.93     | 96         | 609       |      |
| 2 | 1001863 | 120  | 157       | 2760    | 219.28         | 241.53     | 27         | 609       |      |
| 3 | 1000867 | 38   | 24        | 2760    | 242.53         | 242.53     | 204        | 614       |      |
| 4 | 1001080 | 131  | 83        | 2760    | 260.93         | 260.93     | 107        | 659       |      |

In [19]:

```
df.columns
```

Out[19]:

```
Index(['id', 'week', 'center_id', 'meal_id', 'checkout_price', 'base_price',  
      'num_orders', 'city_code', 'region_code', 'center_type', 'op_area',  
      'category', 'cuisine'],  
      dtype='object')
```

In [20]:

```
df['category'].unique()
```

Out[20]:

```
array(['Other Snacks', 'Pasta', 'Desert', 'Rice Bowl', 'Beverages',  
      'Sandwich', 'Pizza', 'Salad', 'Fish', 'Seafood', 'Biryani',  
      'Extras', 'Starters', 'Soup'], dtype=object)
```

In [21]:

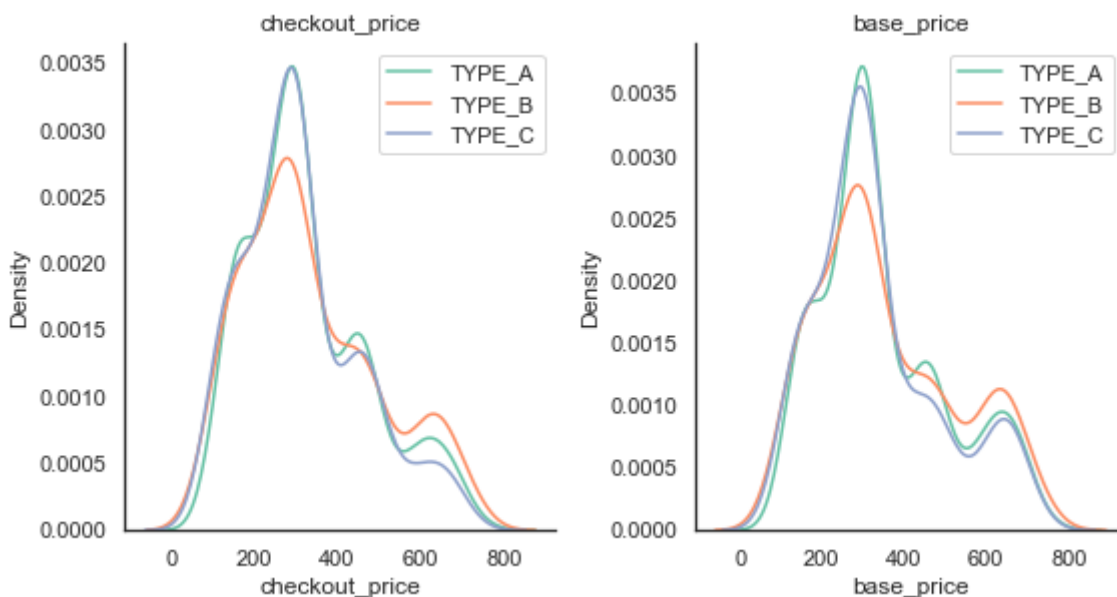
```
df["center_type"].unique()
```

Out[21]:

```
array(['TYPE_A', 'TYPE_B', 'TYPE_C'], dtype=object)
```

In [22]:

```
TYPE_A = df[df["center_type"] == "TYPE_A"]  
TYPE_B = df[df["center_type"] == "TYPE_B"]  
TYPE_C = df[df["center_type"] == "TYPE_C"]  
  
plt.figure(figsize = (16,16))  
for ax, col in enumerate(df.columns[4:6]):  
    plt.subplot(4,4, ax + 1)# here 4,4 will result 4 graphs  
    plt.title(col)  
    sns.kdeplot(x = TYPE_A[col], label = "TYPE_A")  
    sns.kdeplot(x = TYPE_B[col], label = "TYPE_B")  
    sns.kdeplot(x = TYPE_C[col], label = "TYPE_C")  
  
    plt.legend()  
plt.tight_layout()
```



In [23]:

```
df["cuisine"].unique()
```

Out[23]:

```
array(['Thai', 'Italian', 'Indian', 'Continental'], dtype=object)
```

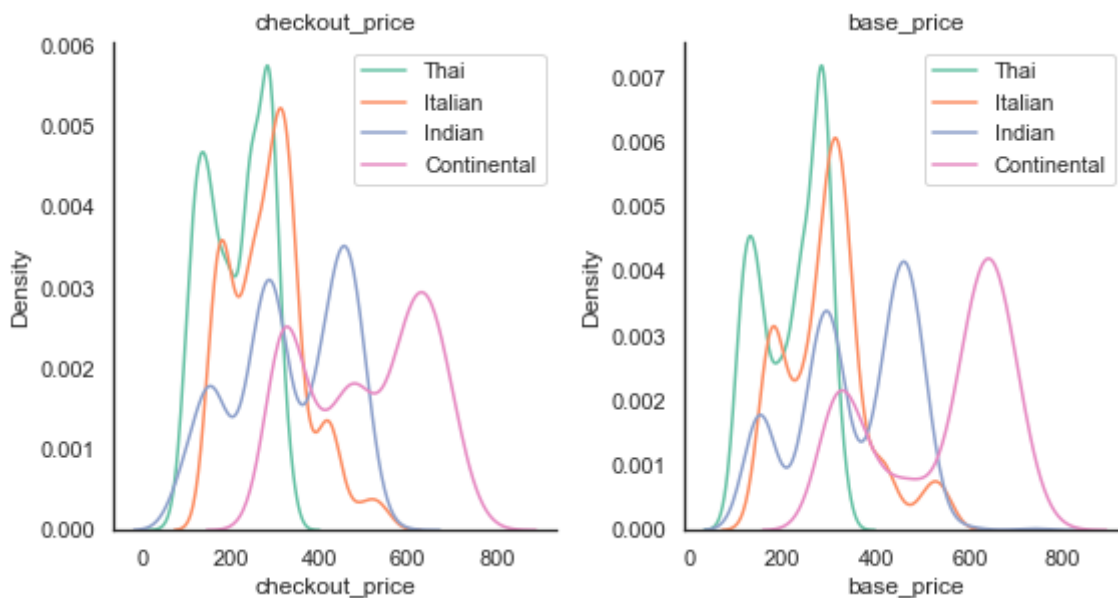


In [24]:

```
A = df[df["cuisine"] == "Thai"]
B = df[df["cuisine"] == "Italian"]
C = df[df["cuisine"] == "Indian"]
D = df[df["cuisine"] == "Continental"]

plt.figure(figsize = (16,16))
for ax, col in enumerate(df.columns[4:6]):
    plt.subplot(4,4, ax + 1)# here 4,4 will result 4 graphs
    plt.title(col)
    sns.kdeplot(x = A[col], label = "Thai")
    sns.kdeplot(x = B[col], label = "Italian")
    sns.kdeplot(x = C[col], label = "Indian")
    sns.kdeplot(x = D[col], label = "Continental")

plt.legend()
plt.tight_layout()
```

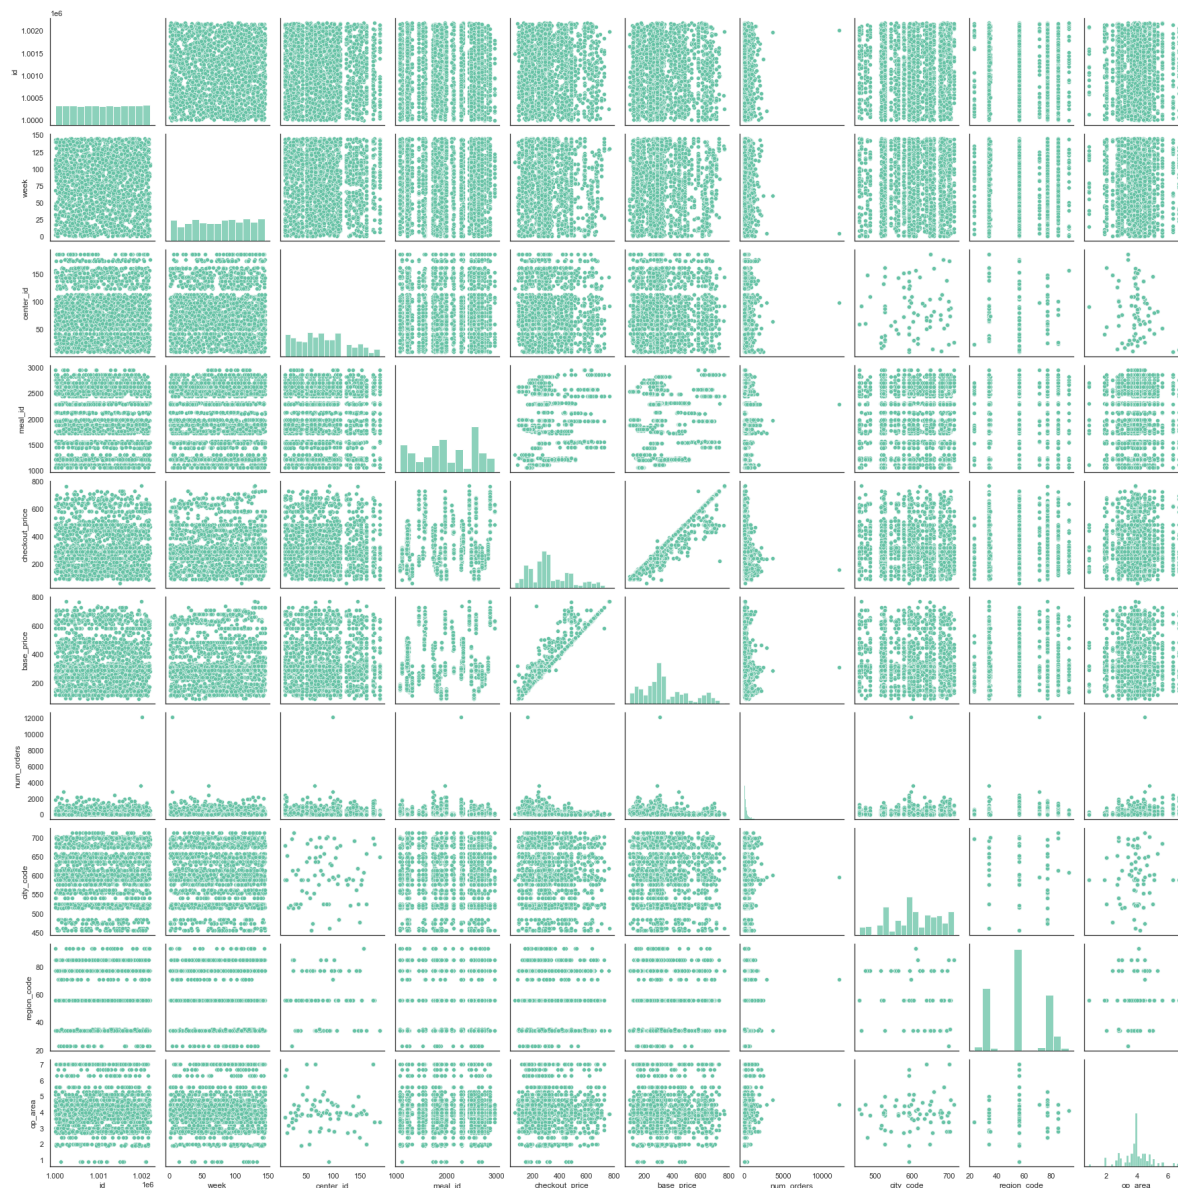


In [25]:

```
sns.pairplot(df)
```

Out[25]:

&lt;seaborn.axisgrid.PairGrid at 0x2120f752ac0&gt;

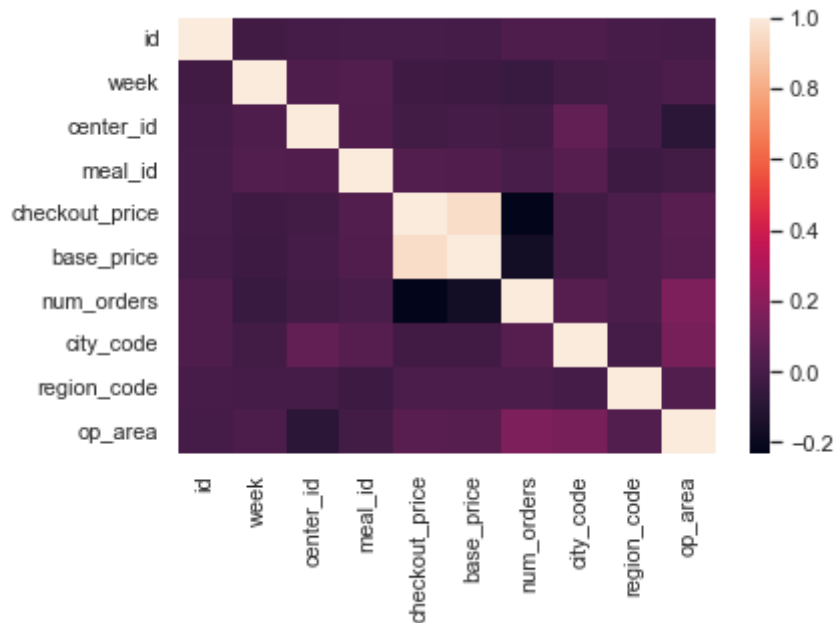


In [26]:

```
corr = df.corr()  
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns)
```

Out[26]:

<AxesSubplot:>



In [28]:

```
from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()  
df['cuisine']=le.fit_transform(df['cuisine'])  
df['center_type']=le.fit_transform(df['center_type'])  
df['category']=le.fit_transform(df['category'])
```

In [29]:

```
df
```

Out[29]:

|      | id      | week | center_id | meal_id | checkout_price | base_price | num_orders | city_code | i |
|------|---------|------|-----------|---------|----------------|------------|------------|-----------|---|
| 0    | 1000000 | 3    | 157       | 2760    | 233.83         | 231.83     | 149        | 609       |   |
| 1    | 1001086 | 144  | 157       | 2760    | 184.36         | 261.93     | 96         | 609       |   |
| 2    | 1001863 | 120  | 157       | 2760    | 219.28         | 241.53     | 27         | 609       |   |
| 3    | 1000867 | 38   | 24        | 2760    | 242.53         | 242.53     | 204        | 614       |   |
| 4    | 1001080 | 131  | 83        | 2760    | 260.93         | 260.93     | 107        | 659       |   |
| ...  | ...     | ...  | ...       | ...     | ...            | ...        | ...        | ...       |   |
| 1994 | 1001791 | 145  | 143       | 2104    | 582.03         | 581.03     | 69         | 562       |   |
| 1995 | 1001295 | 141  | 74        | 2104    | 582.03         | 581.03     | 14         | 702       |   |
| 1996 | 1000382 | 64   | 61        | 2104    | 629.53         | 631.53     | 41         | 473       |   |
| 1997 | 1001089 | 114  | 61        | 2104    | 588.79         | 590.79     | 82         | 473       |   |
| 1998 | 1002144 | 140  | 113       | 2104    | 484.03         | 630.53     | 28         | 680       |   |

1999 rows × 13 columns



In [30]:

```
df=df.drop(columns='id')
```

In [31]:

df

Out[31]:

|      | week | center_id | meal_id | checkout_price | base_price | num_orders | city_code | region_code |
|------|------|-----------|---------|----------------|------------|------------|-----------|-------------|
| 0    | 3    | 157       | 2760    | 233.83         | 231.83     | 149        | 609       | 9           |
| 1    | 144  | 157       | 2760    | 184.36         | 261.93     | 96         | 609       | 9           |
| 2    | 120  | 157       | 2760    | 219.28         | 241.53     | 27         | 609       | 9           |
| 3    | 38   | 24        | 2760    | 242.53         | 242.53     | 204        | 614       | 8           |
| 4    | 131  | 83        | 2760    | 260.93         | 260.93     | 107        | 659       | 7           |
| ...  | ...  | ...       | ...     | ...            | ...        | ...        | ...       | ...         |
| 1994 | 145  | 143       | 2104    | 582.03         | 581.03     | 69         | 562       | 7           |
| 1995 | 141  | 74        | 2104    | 582.03         | 581.03     | 14         | 702       | 9           |
| 1996 | 64   | 61        | 2104    | 629.53         | 631.53     | 41         | 473       | 7           |
| 1997 | 114  | 61        | 2104    | 588.79         | 590.79     | 82         | 473       | 7           |
| 1998 | 140  | 113       | 2104    | 484.03         | 630.53     | 28         | 680       | 7           |

1999 rows × 12 columns

=

In [32]:

```

predictors = [ 'week', 'center_id', 'meal_id', 'checkout_price', 'base_price', 'city_code',
outcome = 'num_orders'

X = df[predictors]
y = df[outcome]

# partition data
train_X, test_X, train_y, test_y = train_test_split(X, y, test_size=0.2, random_state=42)

```

In [33]:

```

print(train_X.shape)
print(test_X.shape)
print(train_y.shape)
print(test_y.shape)

```

```

(1599, 11)
(400, 11)
(1599,)
(400,)

```

In [34]:

```
rf = RandomForestRegressor()  
rf.fit(train_X,train_y)
```

Out[34]:

RandomForestRegressor()

In [35]:

```
# training  
regressionSummary(train_y, rf.predict(train_X))  
# validation  
regressionSummary(test_y, rf.predict(test_X))
```

Regression statistics

```
                Mean Error (ME) : -4.3113  
    Root Mean Squared Error (RMSE) : 139.6324  
        Mean Absolute Error (MAE) : 50.6919  
    Mean Percentage Error (MPE) : -29.4344  
Mean Absolute Percentage Error (MAPE) : 38.0004
```

Regression statistics

```
                Mean Error (ME) : -32.0695  
    Root Mean Squared Error (RMSE) : 234.1408  
        Mean Absolute Error (MAE) : 135.8930  
    Mean Percentage Error (MPE) : -95.6834  
Mean Absolute Percentage Error (MAPE) : 116.0958
```

In [36]:

```
MSE=mean_squared_error(train_y,rf.predict(train_X))  
MSE
```

Out[36]:

19497.2014395247

In [37]:

```
r2=r2_score(train_y,rf.predict(train_X))  
r2
```

Out[37]:

0.8969755105214148

In [38]:

```
lr = LinearRegression()  
lr.fit(train_X,train_y)
```

Out[38]:

LinearRegression()

In [39]:

```
# training
regressionSummary(train_y, lr.predict(train_X))
# validation
regressionSummary(test_y, lr.predict(test_X))
```

Regression statistics

```
Mean Error (ME) : 0.0000
Root Mean Squared Error (RMSE) : 404.2196
Mean Absolute Error (MAE) : 195.1953
Mean Percentage Error (MPE) : -160.8192
Mean Absolute Percentage Error (MAPE) : 219.4809
```

Regression statistics

```
Mean Error (ME) : 3.1304
Root Mean Squared Error (RMSE) : 280.7624
Mean Absolute Error (MAE) : 185.3807
Mean Percentage Error (MPE) : -142.7065
Mean Absolute Percentage Error (MAPE) : 214.3767
```

In [40]:

```
MSE=mean_squared_error(train_y,lr.predict(train_X))
MSE
```

Out[40]:

163393.50022405767

In [41]:

```
r2=r2_score(train_y,lr.predict(train_X))
r2
```

Out[41]:

0.13661804249620557

In [42]:

```
import xgboost as xgb
xgb=xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.3, learning_rate = 0.1,
                      max_depth = 5, alpha = 10, n_estimators = 10)
xgb.fit(train_X,train_y)
```

[23:48:22] WARNING: /workspace/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Out[42]:

XGBRegressor(alpha=10, colsample\_bytree=0.3, max\_depth=5, n\_estimators=10)

In [43]:

```
# training
regressionSummary(train_y, xgb.predict(train_X))
# validation
regressionSummary(test_y, xgb.predict(test_X))
```

Regression statistics

```
                Mean Error (ME) : 93.3079
    Root Mean Squared Error (RMSE) : 359.5107
                Mean Absolute Error (MAE) : 156.0619
                Mean Percentage Error (MPE) : -74.4509
Mean Absolute Percentage Error (MAPE) : 117.1119
```

Regression statistics

```
                Mean Error (ME) : 90.4496
    Root Mean Squared Error (RMSE) : 291.5579
                Mean Absolute Error (MAE) : 160.9225
                Mean Percentage Error (MPE) : -80.4611
Mean Absolute Percentage Error (MAPE) : 124.5857
```

In [44]:

```
MSE=mean_squared_error(train_y,xgb.predict(train_X))
MSE
```

Out[44]:

129247.90955415349

In [45]:

```
r2=r2_score(train_y,xgb.predict(train_X))
r2
```

Out[45]:

0.31704558014169937

In [46]:

```
svm = SVR()
svm.fit(train_X,train_y)
```

Out[46]:

SVR()



In [47]:

```
# training
regressionSummary(train_y, svm.predict(train_X))
# validation
regressionSummary(test_y, svm.predict(test_X))
```

Regression statistics

```
                Mean Error (ME) : 115.6317
      Root Mean Squared Error (RMSE) : 448.8388
                Mean Absolute Error (MAE) : 193.9638
                Mean Percentage Error (MPE) : -106.4076
Mean Absolute Percentage Error (MAPE) : 158.1633
```

Regression statistics

```
                Mean Error (ME) : 113.3291
      Root Mean Squared Error (RMSE) : 331.4003
                Mean Absolute Error (MAE) : 189.7209
                Mean Percentage Error (MPE) : -106.0646
Mean Absolute Percentage Error (MAPE) : 158.5903
```

In [48]:

```
MSE=mean_squared_error(train_y,svm.predict(train_X))
MSE
```

Out[48]:

201456.23876258524

In [49]:

```
r2=r2_score(train_y,svm.predict(train_X))
r2
```

Out[49]:

-0.06450796106137302

In [50]:

```
print(rf.get_params())
```

```
{'bootstrap': True, 'ccp_alpha': 0.0, 'criterion': 'squared_error', 'max_depth': None, 'max_features': 'auto', 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'n_jobs': None, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}
```

In [62]:

```
from sklearn.model_selection import GridSearchCV

param_grid = {
    'bootstrap': [True],
    'max_depth': [80, 90, 100, 110],
    'max_features': [2, 3],
    'min_samples_leaf': [3, 4, 5],
    'min_samples_split': [8, 10, 12],
    'n_estimators': [100, 200, 300, 1000]
}

# Create a based model
rf = RandomForestRegressor()
# Instantiate the grid search model
grid_search = GridSearchCV(estimator = rf, param_grid = param_grid,
                           cv = 3, n_jobs = -1, verbose = 2)

grid_search.fit(train_X, train_y)
grid_search.best_params_
{'bootstrap': True,
 'max_depth': 80,
 'max_features': 3,
 'min_samples_leaf': 5,
 'min_samples_split': 12,
 'n_estimators': 100}
best_grid = grid_search.best_estimator_
grid_accuracy = evaluate(best_grid, test_X, test_y)
print('Improvement of {:.2f}%'.format( 100 * (grid_accuracy - base_accuracy) / base_accu
```

Fitting 3 folds for each of 288 candidates, totalling 864 fits

Model Performance

Average Error: 131.8716 degrees.

Accuracy = 23.35%.

Improvement of 45.09%.